

Computing Pure Nash Equilibria in Symmetric Action Graph Games

Albert Xin Jiang Kevin Leyton-Brown
Department of Computer Science
University of British Columbia
{jiang;kevinlb}@cs.ubc.ca


INFORMS: October 14, 2008

Outline


- 1 Action Graph Games
- 2 Computing Pure Nash Equilibria
- 3 Computing Pure Equilibria in Symmetric AGGs
- 4 Algorithm
- 5 Conclusions

Example: Location Game

- each of n agents wants to open a business
- actions: choosing locations
- utility: depends on
 - the location chosen
 - number of agents choosing the same location
 - numbers of agents choosing each of the adjacent locations



Game on a graph



- This can be modeled as a game played on a directed graph:
 - each player has a token to put on one of the nodes;
 - each player's utility depends on:
 - the node chosen
 - configuration of tokens over neighboring nodes
- Action Graph Games (Bhat & Leyton-Brown 2004, Jiang & Leyton-Brown 2006)
 - fully expressive, compact representation of games
 - exploits anonymity, context specific independence

Definitions

Definition (action graph)

An **action graph** is a tuple (\mathcal{A}, E) , where \mathcal{A} is a set of nodes corresponding to *distinct actions* and E is a set of directed edges.

- Each agent i 's set of available actions: $A_i \subseteq \mathcal{A}$
- Neighborhood of node α : $\nu(\alpha) \equiv \{\alpha' \in \mathcal{A} \mid (\alpha', \alpha) \in E\}$

Definitions

Definition (action graph)

An **action graph** is a tuple (\mathcal{A}, E) , where \mathcal{A} is a set of nodes corresponding to *distinct actions* and E is a set of directed edges.

- Each agent i 's set of available actions: $A_i \subseteq \mathcal{A}$
- Neighborhood of node α : $\nu(\alpha) \equiv \{\alpha' \in \mathcal{A} \mid (\alpha', \alpha) \in E\}$

Definition (configuration)

A **configuration** c is an $|\mathcal{A}|$ -tuple of integers $(c[\alpha])_{\alpha \in \mathcal{A}}$. $c[\alpha]$ is the number of agents who chose the action $\alpha \in \mathcal{A}$. For a subset of actions $X \subset \mathcal{A}$, let $c[X]$ denote the restriction of c to X . Let $C[X]$ denote the set of restricted configurations over X .

Action Graph Games

Definition (Action Graph Game (AGG))

An **action graph game** Γ is a tuple $\langle N, (A_i)_{i \in N}, G, u \rangle$ where

- N is the set of agents
- A_i is agent i 's set of actions
- $G = (\mathcal{A}, E)$ is the action graph, where $\mathcal{A} = \bigcup_{i \in N} A_i$ is the set of distinct actions
- $u = (u^\alpha)_{\alpha \in \mathcal{A}}$, where $u^\alpha : C[\nu(\alpha)] \mapsto \mathbb{R}$

Action Graph Games

Definition (Action Graph Game (AGG))

An **action graph game** Γ is a tuple $\langle N, (A_i)_{i \in N}, G, u \rangle$ where

- N is the set of agents
- A_i is agent i 's set of actions
- $G = (\mathcal{A}, E)$ is the action graph, where $\mathcal{A} = \bigcup_{i \in N} A_i$ is the set of distinct actions
- $u = (u^\alpha)_{\alpha \in \mathcal{A}}$, where $u^\alpha : C[\nu(\alpha)] \mapsto \mathbb{R}$

Definition (symmetric AGG)

An AGG is **symmetric** if all players have identical action sets, i.e. if $A_i = \mathcal{A}$ for all i .

AGG Properties

- AGGs are **fully expressive**
- Symmetric AGGs can represent **arbitrary symmetric games**
- **Representation size $\|I\|$ is polynomial** if the in-degree \mathcal{I} of G is bounded by a constant
- **Any graphical game** (Kearns, Littman & Singh 2001) can be encoded as an AGG of the same space complexity.
- AGG can be **exponentially smaller** than the equivalent graphical game & normal form representations.

Outline

- 1 Action Graph Games
- 2 Computing Pure Nash Equilibria**
- 3 Computing Pure Equilibria in Symmetric AGGs
- 4 Algorithm
- 5 Conclusions

Pure Nash Equilibria

Action profile: $\mathbf{a} = (a_1, \dots, a_n)$

Definition (pure Nash equilibrium)

An action profile \mathbf{a} is a **pure Nash equilibrium** of the game Γ if for all $i \in N$, a_i is a best response to a_{-i} (i.e. for all $a'_i \in A_i$, $u_i(a_i, a_{-i}) \geq u_i(a'_i, a_{-i})$).

- not guaranteed to exist
- often more interesting than mixed Nash equilibria

Complexity of Finding Pure Equilibria

Checking every action profile:

- linear time in normal form size
- worst-case **exponential time** in AGG size

Complexity of Finding Pure Equilibria

Checking every action profile:

- linear time in normal form size
- worst-case **exponential time** in AGG size

Consider the restriction to symmetric AGGs.

Theorem (Conitzer, personal communication; also proven independently in (Daskalakis *et al.* 2008))

The problem of determining whether a pure Nash equilibrium exists in a symmetric AGG is NP-complete, even when the in-degree of the action graph is at most 3.

Our Contribution

We provide an algorithm that is tractable for **symmetric AGGs with bounded treewidth**

- the algorithm can also be applied to other settings

Specifically, we propose a **dynamic programming** approach:

- partition action graph into **subgraphs** (via tree decomposition)
- construct equilibria of the game from equilibria of games played on subgraphs

Our Contribution

We provide an algorithm that is tractable for **symmetric AGGs with bounded treewidth**

- the algorithm can also be applied to other settings

Specifically, we propose a **dynamic programming** approach:

- partition action graph into **subgraphs** (via tree decomposition)
- construct equilibria of the game from equilibria of games played on subgraphs

Related Work:

- finding pure equilibria in **graphical games**
 - (Gottlob, Greco, & Scarcello 2003) and (Daskalakis & Papadimitriou 2006)
- finding pure equilibria in **simple congestion games**
 - (Ieong, McGrew, Nudelman, Shoham, & Sun 2005)


Outline

- 1 Action Graph Games
- 2 Computing Pure Nash Equilibria
- 3 Computing Pure Equilibria in Symmetric AGGs**
- 4 Algorithm
- 5 Conclusions

Restricted Game

To derive an algorithm that builds up from partial solutions, we must define the concept of a **restricted game**


- game played by a subset of players: $n' \leq n$
- actions restricted to $R \subseteq \mathcal{A}$
- utility functions same as in original AGG
 - need to specify configuration of neighboring nodes not in R



- *restricted game* $\Gamma(n', R, c[\nu(R)])$

Partial Solution

We want to use equilibria of restricted games as building blocks




Definition (partial solution)

A **partial solution** on a restricted game $\Gamma(n', X, c[\nu(X)])$ is a configuration $c[X \cup \nu(X)]$ such that $c[X]$ is a pure NE of Γ .

Extending partial solutions

- **Problem:** combining two partial solutions on two non-overlapping restricted games does not necessarily produce an equilibrium of the combined game
 - configurations may be **inconsistent**, or
 - player might **profitably deviate** from playing in one restricted game to another
- keeping all partial solutions: impractical as sizes of restricted games grow
- we would like **sufficient statistics** that summarize partial solutions as compactly as possible

Sufficient statistic



Sufficient Statistic: a tuple consisting of

1. **configuration** over
 - outside neighbours: $\nu(X)$
 - inside nodes that are neighbors of outside nodes: $\nu(\overline{X})$
2. **number of agents** playing in X
3. U_w , **utility of the worst-off player** in $X \setminus \nu(\overline{X})$.
4. U_b , **best utility an outside player** can get by playing in $X \setminus \nu(\overline{X})$.

Number of distinct tuples: **polynomial** for action graphs of bounded treewidth

Combining sufficient statistics

Given two sets of such tuples, summarizing partial solutions on $X, Y \subset \mathcal{A}$, we can compute the set of sufficient statistics for the **combined** restricted game $X \cup Y$

- start with all **consistent** configurations
 - analogous to database join of the two sets of tuples
- discard those with **profitable $X \rightarrow Y$ deviations** (& vice versa)
 - easy: discard when U_w from X is worse than U_b from Y
 - trickier: checking deviations from $X \cap \nu(Y)$ to $\nu(\bar{Y})$
 - utilities in $\nu(\bar{Y})$ change when $c[\nu(Y)]$ changes, so checking these deviations is more costly
 - solution: augment our sufficient statistics to keep track of the **configuration of the neighborhood** of $\nu(\bar{Y})$, in order to compute these utilities on the fly
 - luckily, for graphs of bounded treewidth, this implies storing a small amount of additional information
- overall: all profitable deviations can be discarded efficiently

Outline

- 1 Action Graph Games
- 2 Computing Pure Nash Equilibria
- 3 Computing Pure Equilibria in Symmetric AGGs
- 4 Algorithm**
- 5 Conclusions

Algorithm

- 1 Construct the **primal graph** of the action graph.
- 2 Build a **tree decomposition** of this primal graph.
- 3 Partition the AGG according to the tree decomposition.
- 4 Find all sufficient statistics¹ corresponding to **partial solutions of games restricted to each partition**.
- 5 Working up the tree, **combine adjacent nodes** together.
- 6 When **root is reached**, return whether the game has a PSNE.

¹Augment sufficient statistics to include configurations over additional actions that belong to the decomposition's tree node that is closest to the root.

Algorithm

- 1 Construct the **primal graph** of the action graph.
- 2 Build a **tree decomposition** of this primal graph.
- 3 Partition the AGG according to the tree decomposition.
- 4 Find all sufficient statistics¹ corresponding to **partial solutions of games restricted to each partition**.
- 5 Working up the tree, **combine adjacent nodes** together.
- 6 When **root is reached**, return whether the game has a PSNE.


Theorem

*For symmetric AGGs with bounded treewidth, our algorithm determines existence of pure Nash equilibria in **polynomial time**.*

Recover a PSNE from the SS's: downwards pass on the tree

¹Augment sufficient statistics to include configurations over additional actions that belong to the decomposition's tree node that is closest to the root.


An Example




- Two players
- Utility functions:
 - start with payoff of 0
 - +1 reward if playing action F or D
 - -2 penalty if another player selected an action with an incoming edge
 - For C , this means a neighboring action (since C does not have a self-edge)
 - Otherwise, this means the same or a neighboring action
- Pure Nash equilibria:
 - One player chooses D , the other chooses F
 - Both players choose C

1. Construct Primal Graph

Action graph:




Primal graph: make each neighborhood a clique



2. Construct Tree Decomposition

A tree where each node is labeled with one or more nodes from the primal graph, where

- every label is used **at least once**
- for every edge in the primal graph from α_1 to α_2 , there is a node in the tree **labeled with both** α_1 and α_2
- if a label occurs in two nodes x_1, x_2 in the tree, it also **occurs on all paths** between x_1 and x_2 .




If treewidth of the AGG is bounded by a constant, the primal graph's tree decomposition can be **computed in polynomial time**.



3. Partition the AGG According to the Tree Decomposition

By construction: for each node α in the action graph, there always exists a tree node in the decomposition of the primal graph that contains α and its neighbors in the action graph.

The tree decomposition therefore induces the following partition on the AGG:



4. Compute Sufficient Statistics for Partial Solutions on Each Partition



For restricted game on $\{C\}$:


n'	$c[B, C, D, F]$	$U_w(\emptyset)$	$U_b(\emptyset)$
0	0,0,0,0	∞	$-\infty$
0	1,0,0,0	∞	$-\infty$
...	...	∞	$-\infty$
1	0,1,0,0	∞	$-\infty$
1	1,1,0,0	∞	$-\infty$
...	...	∞	$-\infty$
2	0,2,0,0	∞	$-\infty$

For restricted game on $\{F, G\}$:

n'	$c[C, F, G]$	$U_w(G)$	$U_b(G)$
0	0,0,0	∞	0
0	1,0,0	∞	0
0	2,0,0	∞	0
1	0,1,0	∞	-2
1	1,0,1	0	-2
2	0,1,1	-2	$-\infty$

5. Working up the Tree, Combine Restricted Games

Combine restricted games in bottom-up order: from leaves to root.




Combine $\{C\}$ and $\{F, G\}$ to create table for restricted game on $\{C, F, G\}$:

n'	$c[B, C, D, F]$	$U_w(G)$	$U_b(G)$
0	0,0,0,0	∞	0
0	1,0,0,0	∞	0
...	...	∞	0
1	0,0,0,1	∞	-2
1	1,0,0,1	∞	-2
1	0,0,1,1	∞	-2
2	0,1,0,0	0	$-\infty$
2	0,2,0,0	∞	$-\infty$

5. Working up the Tree, Combine Restricted Games

Combine restricted games in bottom-up order: from leaves to root.





Combine $\{D, E\}$ and $\{C, F, G\}$ to create table for $\{C, D, E, F, G\}$:

n'	$c[B, C, D, F]$	$U_w(E, G)$	$U_b(E, G)$
0	0,0,0,0	∞	0
0	1,0,0,0	∞	0
0	2,0,0,0	∞	0
1	0,0,1,0	∞	0
1	1,0,1,0	∞	0
1	0,0,0,1	∞	0
1	1,0,0,1	∞	0
2	0,0,1,1	∞	$-\infty$
2	0,2,0,0	∞	$-\infty$

5. Working up the Tree, Combine Restricted Games

Combine restricted games in bottom-up order: from leaves to root.



Combine $\{A, B\}$ and $\{C, D, E, F, G\}$:

n'	$c[A, B, C]$	$U_w(D, E, F, G)$	$U_b(D, E, F, G)$
2	0,0,0	1	$-\infty$
2	0,0,2	∞	$-\infty$

6. Top-Down Pass to Compute PNSE

n'	$c[A, B, C]$	$U_w(D, E, F, G)$	$U_b(D, E, F, G)$
2	0,0,0	1	$-\infty$
2	0,0,2	∞	$-\infty$

To compute a PSNE, start from the root and work down. At each node, pick a row from the table of sufficient statistics that is consistent with earlier picks.

- If we start with row 1, we select an equilibrium in which one player chooses D , one player chooses F
- If we start with row 2, we select an equilibrium in which both players choose C

Outline

- 1 Action Graph Games
- 2 Computing Pure Nash Equilibria
- 3 Computing Pure Equilibria in Symmetric AGGs
- 4 Algorithm
- 5 Conclusions**

Conclusions & Beyond Symmetric AGGs

- dynamic programming approach for computing pure equilibria in AGGs
- poly-time algorithm for symmetric AGGs with bounded **treewidth**
- our approach can be extended to general AGGs
 - different set of sufficient statistics
 - when the game is k -symmetric (i.e. has k distinct action sets), use k -configuration (k -tuple of configurations, one for each equivalence class of players), and similarly use k -tuples of U_w, U_b
 - for subgraphs in which only k' of the k classes of players participate, only need to keep track of the sufficient statistics for those k' classes.
 - related algorithms for graphical games (Daskalakis & Papadimitriou 2006) and simple congestion games (leong et al 2005) become special cases of our approach